

CS-GY 6513 Big Data Final Project

ROBERT RONAN, rr3749

SHENG TONG, st4048

JERRY LEE, jl11517

Github repo: <https://github.com/darkcb109/DOB-Job-Application-Cleaning>

1 INTRODUCTION

For this project we profiled, analyzed, and cleaned the New York City Department of Buildings “Job Applications Filings” dataset, which is publicly available on data.cityofnewyork.us/. As more and more data is increasingly collected, and as government agencies continue to publish more and more public data sets, data cleaning will become increasingly important. Data sets frequently contain erroneous values, and data are stored in multiple, incompatible formats. When data sets that have not been cleaned are joined, this problem can become even worse.

In this project, we cleaned the DOB Job Application Filings dataset, applied the techniques we used on that data set to ten similar data sets, and measured the results. We then used precision and recall on a subset of the data to measure the accuracy of our techniques. All of our code is written in python using tools like Pandas and VIDA-NYU’s OpenClean library. We will take a look at some of the problems we encountered and some methods to try and ease those problems. Then, we will discuss our findings and other problems found that were a bit tricky to solve.

The objective of our project is to find and fix most of the problems in a large dataset and as well as handle an even larger dataset. We take a look at many of the problems addressed in the DOB Job Application Filing dataset and try to broaden the scope such as names and numbers. Then we find a fix to these and apply them to multiple datasets. This cycle continues until the entire dataset is covered.

2 PROBLEM FORMULATION

First, we took a look at some columns and noticed that some column names were vague and inconsistent with the rest. Our plan was to rename misspellings, capitalization inconsistencies, as well as how they use certain names and symbols such as “No.”, “”, “Number” which can all be changed to ‘Number’. Other naming issues we found were vague columns where “City” was referring to the owner’s house rather than the city where the job filling is held.

Next, we looked at some fields with some simple profiling and separated some problems to look at based on the type of field (numeric, monetary, datetime, phone numbers, binary data, etc) and changed them to match consistency as well as correct some improper fields. The biggest problem we discovered were the inconsistent uses of None, NaN, and blank space. There are also other problems involving incorrect data such as zip codes outside NYC, misspelling of

Authors’ netid: Robert Ronan, rr3749; Sheng Tong, st4048; Jerry Lee, jl11517.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2021 Association for Computing Machinery.

Manuscript submitted to ACM

Manuscript submitted to ACM

names, and problems with inconsistent name values (NYC, New York, New York City). One problem that we discovered that was harder to fix was the coherency between certain values such as assigned dates being later than approved dates, however since the fraction of jobs that had this was very minimal so we left it as is. We used kNN clusters to find some outliers to some names, cities, and dates. We also took a quick glance at the GIS fields such as latitude and longitude and noticed that there were 0 fields and looking further, found out that the entire data was filled with incorrect information and we decided to remove the entire row.

After looking at every field and updating our values, we tried to broaden our scope so it would work on different but similar datasets. We took a look at other Department of Buildings datasets since they had many similar fields. A problem that our original data did that we didn't account for was that many of the fields are hard-coded into our data cleaning. We tried to complement this issue by making a very general model of our data cleaning. Although there were other fields that couldn't be generalized and had to be manually cleaned. One of these is the usage of clusters on extremely large datasets which will tend to use up way too much memory usage. Some fields are also too ambiguous which we wouldn't be able to tell if it is a certain type of dataset (ex. 10 numbers can mean both phone and some form of identification).

3 RELATED WORK

Since we used Pandas and OpenClean to help us with data cleaning, we followed the OpenClean examples provided in the OpenClean repository on data profiling and data cleaning. Although we didn't use many of the OpenClean functions since there were many issues running the DOB Job Applications Filings dataset in most of our environment, we used Pandas to clean and profile easier fields. However, clustering and soundex was much easier to implement so we used the OpenClean clustering functions to determine outliers and cleaned the data. Link to OpenClean repository: <https://github.com/VIDA-NYU/openclean-core>

4 METHODS, ARCHITECTURE AND DESIGN

We utilized a number of distinct methods to profile and clean the DOB Job Application Filings data set. Below, we will discuss different types of errors and inconsistencies we found, and how we attempted to address them.

4.1 Column Names

We identified a number of issues with the original column names. Some column names were missing spaces between words, some had multiple spaces. Some column names used underscores to separate words, and some used spaces. Different column names used entirely different conventions for capitalization and abbreviations. For example, the word "number" was sometimes written out entirely, sometimes written as "No." and sometimes as a . Words in column names were occasionally abbreviated, whereas other column names which were significantly longer contained no abbreviations. Finally, some column names were ambiguous to the point of causing confusion. The column "Paid", did not actually indicate a boolean value, but was the date that payment was rendered on.

To address these issues, we settled on a uniform naming convention for the column names. We replaced all underscores with spaces, made columns title case except for all capital abbreviations. We changed all versions of "number" to the word "number"; and we appended "Date" to datetime column names that did not contain it.

4.2 Numerical Columns

For numerical columns, our goal was ultimately to convert them into integers as floats. To this end, we removed non-numeric values and dollar signs. Doing this allowed us to then check the distribution of the column values more easily, and to spot outliers. In particular, checking the minimum and maximum values of the columns often identified obvious outliers and errors. When looking at monetary values, we discovered a number of instances in which the amount of money listed was astronomically large, or was negative in cases where that would not have made sense. We also looked at columns containing the proposed number of feet and stories for buildings, and we were able to spot erroneously large values by comparing these against the tallest buildings in New York City.

Unfortunately, while these methods allow us to address issues in our numerical columns, many of them are context specific and difficult to automate. For instance, none of our numerical values could be negative, but this is clearly not true in general for arbitrary numerical columns. Examining outlying values in these columns allows us to identify a number of problems, but also requires ‘expert knowledge’ of the columns, which is not always present.

4.3 Boolean/Binary Columns

We identified a number of columns that were implicitly boolean, and which were more usefully represented with boolean values. Often these column names were questions in the form of statements like, “pre-filed”, or “landmark.” Frequently they contained values of “Yes/No” or “Y/N”, but often a significant number of NaNs. We identified numerous columns that appear to have been checkboxes at some point because their only values were “X” and NaN.

For these columns, we converted values indicating yes, including “X” to True, and other values to False, including NaNs. Some of these columns also contained a very small number of erroneous values, like one instance of the letter ‘H.’ However, There were columns that looked very close to binary, but did have legitimate third or fourth options, so our approach was somewhat column specific.

4.4 Phone Numbers

Columns containing phone numbers were riddled with errors. There was no consistent format for the phone numbers (some had dashes, some spaces, some parentheses). Frequently, the listed phone numbers were missing digits or had too many.

For the phone numbers, we stripped all non-numeric characters, and dropped all numbers that were not ten digits, or that started with a 0 or a 1, which are not valid phone numbers in the United States.

4.5 House Numbers

House numbers were difficult to parse and identify errors in because house numbering is not as uniform as one would hope. Some house numbers were spelled out words, which we removed, but since there was no apartment column, many of the house numbers included either letters or the words “Garage” or “Rear”.

While something like “200 Garage” may not typically be considered a “house number”, our data set was about Department of Buildings construction job applications, and so in our case, it is likely very relevant whether the work is happening on the garage instead of the house, or on the apartment in the rear of the house. Thus, we standardized the formatting by inserting spaces between the numerical component of the house number and any letters.

4.6 New York City Specific Information

Here we examined the associated Council Districts, Latitude and Longitude, and confirmed that all of these values fell in the appropriate range for New York City based data.

4.7 DateTime Data

Primarily the issues we identified with datetime data was that dates were stored seemingly randomly in month/day/year format, and in year-month-day format. After fixing the datetime formatting, we checked values from multiple datetime columns for coherency.

We used Pandas’s built-in `to_datetime` function to convert all the datetimes to the preferred year-month-day format. In a contextual manner, we checked the coherency of various datetime columns by assessing what proportion of them indicated events occurring in the wrong order. For example we looked for data instances where the “Paid Date” was strictly after the “Fully-Paid Date”, or where the job was approved before it was assigned, or where the ‘Pre-Filing Date” was after the “Latest Action Date.” We discovered that in a relatively small number of examples, one of these was the case, but as we lacked the expert knowledge of whether this was technically possible in some way, we did not remove these values.

4.8 Name-Related Data

For columns with name data, like business names, applicant professional titles, first and last names, we used OpenClean’s K-Nearest-Neighbor clustering to identify and rectify outliers, as well as to standardize the formatting. There were more potential formatting problems than we could reasonably list here, but we identified no consistency in capitalization, punctuation, abbreviation, or alliteration. We also identified typos.

Given the breadth of potential errors, we figured our best bet would be to use some type of clustering algorithm to remove formatting issues, and fix misspellings. For this task, we used a very high confidence level that multiple values were in the same cluster before we associated them all with the clustered value. This approach proved quite successful in addressing punctuation, capitalization, and some other issues like abbreviations. However, we did notice multiple clusters that mapped a number of correct, but distinct values to the same name. It is possible that we needed to tweak the settings more. We were also unable to cluster on some columns we had hoped to, due to their size and complexity.

4.9 Cities

For city name data (like (business) owner’s city), we identified issues by using OpenClean’s Soundex implementation alongside an encyclopedia of United States Cities. We identified frequent misspellings of cities, and by also using OpenClean’s functional dependency violations module, we were able to correct a number of incorrect city-state pairs that users had entered.

Many of these errors were corrected by hand, which unfortunately did not scale well to other data sets. However, since all the data sets were New York City based, some corrections still applied.

5 RESULTS

We ran our data over a sample of 50 sample data and measured our precision and recall over these data over what columns that we changed and here are our results.

Dataset	Precision (Original)	Recall (Original)	Precision (Improved)	Recall (Improved)
DOB Cellular Antenna Filings	0.995	0.577	0.863	0.984
Open Restaurant Applications	0.474	0.538	0.995	0.986
DOB NOW: Safety – Facades Compliance Filings	0.996	0.793	0.864	0.979
DOB Permit Issuance	0.990	0.746	0.987	0.984
DOB NOW: Build	0.947	0.904	0.944	0.962
DOB NOW: Electrical Permit Applications	0.986	0.917	0.877	0.903
Historical_DOB_Permit	0.945	0.738	0.977	0.843
Housing New York Units by Building	0.933	0.767	1.0	0.885
DOB NOW: Build – Approved Permits	0.975	0.826	0.845	0.760
Buildings Selected for the Alternative Enforcement Program (AEP)	1.0	0.778	0.935	1.0

The precision and recall differs heavily from each dataset, which can be assumed that our original methods didn't cover all fields. This also means that we know what parts need to be fixed. These 10 new datasets have similar columns, but the data is very different. What we did is to manually look into each of the datasets and do data profiling and then do data cleaning. Here are problems we found during the process:

A majority of our false negatives were from NaN values in non-numeric values. Our decision was to make empty or none for numerical values NaN while names become an empty string, and NaT for time values. However, there were some parts that couldn't be changed without affecting other values. One of these is the names of businesses, and also people's first and last names. In the Job Applications Filings database, we ran into an issue where "NYC Housing Authority" had many different spellings which led to a problem because if we used a stronger classification, it would mess with other businesses which are similar but distinct, so we came to a consensus to use a low eps value.

For datetime data, we have different date formats in different datasets, and sometimes there are different format outliers in those data. We can not simply convert all of them into one specific format, and sometimes there are illegal values which will cause errors when we try to convert string type into datetime type.

Also, there is a serve problem about our method, we need to carefully look into each dataset to decide how to clean the data in detail, in this way we will have high quality of cleaned data, but it requires too much time and effort to perform on every dataset in the NYC Open Data database.

With that in mind, we have improved our method by generalizing our techniques and code, so the code now can automatically find columns to clean, and determine how to clean the data based on our experience of manually cleaning these datasets. Our Improved version sometimes yields slightly worse results, since the improvement we made to our methods was to generalize every method to work on a multitude of datasets rather than hard code some values. This would naturally mean that some of the edge cases would not be covered. For example, for binary/pseudo-binary columns we used to convert several kinds of characters into Y and set all other characters as N, but in some datasets there are outliers like “Yes” or “O” to present Y. And also, it’s even hard to find binary/pseudo-binary columns because we can not find keywords for locating those columns.

And what is worse, sometimes we will clean one column twice in different ways, for instance, we automatically find columns related to “city” and “name” to fix, but then the column “City Name” will be first treated as city and then treated as name which will cause a lot of false positives.

After several changes and improvements to our generalized code, we had it automatically run on all 10 datasets, and we manually evaluated its result in terms of precision and recall. As the chart shown above, we got relatively normal results compared to our original methods, precision decreased for most of the datasets and recalls varied on different datasets.

Precisions are lower than before might because we generalized every method based on our experience as mentioned above, some unwanted changes happened during the cleaning process. As for recalls, because we automatically take columns, sometimes we might take too many unrelated columns. Also, because some of the string data is too long, we can not perform clustering on those data, or we will cause a memory error. So we didn’t perform any clustering in our generalized code, and this also contributes to the recall differences.

6 DISCUSSION

We generalized our code in such a way that we hope it could run on most of the NYC Open Data datasets. To run our cleaning strategy over all datasets would require it to be entirely column-name agnostic, and to infer the data types directly from the columns themselves. We have partially done this; by trying to infer column data types from a combination of non-hardcoded column names and data values.

For instance, we search column names for the word ‘date’, and presume that any column with ‘date’ in its name is a datetime column, which we proceed to clean as if it were. Similarly, we search for columns with ‘number’ in their names to try and determine numerical columns. From those columns, we also search for “building” and “number” or “house” and “number” to find house numbers which we clean in a specific way, as they are more akin to strings than numerical values. We also try to infer when columns are adequately represented as boolean columns, by searching for columns with a small number of distinct values that also contain values like “yes”, or “N”, or True.

Our implementation of this means that both technically and practically, our code be run across all NYC Open Data datasets to clean then, however, because we’ve made our code so much more general, it’s likely clean numerous columns, but with possibly large numbers of false positives and false negatives.

To provide equally broad, but more accurate cleaning across all these datasets, we would likely have to implement, or more effectively utilize, something similar to what Pandas attempts to do in inferring column types when the information is not provided.

7 REFERENCES

[1] DOB Job Application Filings

<https://data.cityofnewyork.us/Housing-Development/DOB-Job-Application-Filings/ic3t-wcy2>

[2] Housing New York Units by Building

<https://data.cityofnewyork.us/Housing-Development/Housing-New-York-Units-by-Building/hg8x-zxpr>

[3] DOB NOW: Build – Job Application Filings

<https://data.cityofnewyork.us/Housing-Development/DOB-NOW-Build-Job-Application-Filings/w9ak-ipjd>

[4] DOB NOW: Electrical Permit Applications

<https://data.cityofnewyork.us/City-Government/DOB-NOW-Electrical-Permit-Applications/dm9a-ab7w>

[5] DOB NOW: Build – Approved Permits

<https://data.cityofnewyork.us/Housing-Development/DOB-NOW-Build-Approved-Permits/rbx6-tga4>

[6] DOB Permit Issuance

<https://data.cityofnewyork.us/Housing-Development/DOB-Permit-Issuance/ipu4-2q9a>

[7] DOB NOW: Safety – Facades Compliance Filings

<https://data.cityofnewyork.us/Housing-Development/DOB-NOW-Safety-Facades-Compliance-Filings/xubg-57si>

[8] Historical DOB Permit Issuance

<https://data.cityofnewyork.us/Housing-Development/Historical-DOB-Permit-Issuance/bty7-2jhb>

[9] Buildings Selected for the Alternative Enforcement Program (AEP) <https://data.cityofnewyork.us/Housing-Development/Buildings-Selected-for-the-Alternative-Enforcement/hcir-3275>

[10] Open Restaurant Applications

<https://data.cityofnewyork.us/Transportation/Open-Restaurant-Applications/pitm-atqc>

[11] DOB Cellular Antenna Filings

<https://data.cityofnewyork.us/Housing-Development/DOB-Cellular-Antenna-Filings/iz2q-9x8d>

[12] OpenClean

<https://github.com/VIDA-NYU/openclean-core>